

2. Relational Model

CSCI 2541 Database Systems & Team Projects

Wood & Chaufournier

Admin Stuff

HW 1 due yesterday

- Test cases should give you idea of your grade

Optional RPS HW due Jan 23

Read our materials/instructions carefully

Read the syllabus for course policies

Watch Slack for ways to **#engage** and ask
#questions

Last week...

Structure that is
independent of
the underlying
file formats

Queries to
flexibly read,
update, and
delete
information

Transactions
that provide
guarantees
about **multi-user**
consistency

Relational Model
Definitions

Constraints and
Relationships

Lab!

...this week.

Data

Let's store some information about professors

- How?

Tables

A Table is a set of rows and columns...

- A column defines an attribute that can have different values
- A row represents related attributes that together represent a data element

Instructor Table

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Course Table

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-101	Intro. to Biology	Biology	4
BIO-301	Genetics	Biology	4
BIO-399	Computational Biology	Biology	3
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3
EE-181	Intro. to Digital Systems	Elec. Eng.	3
FIN-201	Investment Banking	Finance	3
HIS-351	World History	History	3
MU-199	Music Video Production	Music	3
PHY-101	Physical Principles	Physics	4

Tables = Relations

A Relation is a set of tuples and attributes

- Set: an unordered list of unique elements Why?
- Tuple: a sequence of values
- Attribute: a named type with values in a domain

Instructor Relation

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Course Relation

course_id	title	dept_name	credits
BIO-101	Intro. to Biology	Biology	4
BIO-301	Genetics	Biology	4
BIO-399	Computational Biology	Biology	3
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3
EE-181	Intro. to Digital Systems	Elec. Eng.	3
FIN-201	Investment Banking	Finance	3
HIS-351	World History	History	3
MU-199	Music Video Production	Music	3
PHY-101	Physical Principles	Physics	4

Schema

Defines the structure of one or more Relations

- **A₁, A₂, ..., A_n** are attributes
- **R = (A₁, A₂, ..., A_n)** is a relation schema

Example: **instructor = (ID, name, dept_name, salary)**

- A relation instance **r** defined over schema **R** is denoted by **r (R)**.
- The current values of a relation are specified by a table
- An element **t** of relation **r** is called a tuple and is represented by a row in a table

Example DB Schema

STUDENT

Name	Student_number	Class	Major
------	----------------	-------	-------

COURSE

Course_name	Course_number	Credit_hours	Department
-------------	---------------	--------------	------------

PREREQUISITE

Course_number	Prerequisite_number
---------------	---------------------

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
--------------------	---------------	----------	------	------------

GRADE_REPORT

Student_number	Section_identifier	Grade
----------------	--------------------	-------

Relational Model Definitions

A **relation** is a table with columns and rows.

An **attribute** is a named column of a relation.

A **tuple** is a row of a relation.

A **domain** is a set of allowable values for one or more attributes.

The **degree** of a relation is the number of attributes it contains.

The **cardinality** of a relation is the number of tuples it contains.

A **relational database** is a collection of normalized relations with distinct relation names.

Definitions

Degree =

Cardinality =

Help me fill
these in!

The diagram illustrates a database table structure. At the top, four labels point to specific parts of the table: 'attributes (or columns)' points to the column headers ('ID', 'name', 'dept_name', 'salary'); 'tuples (or rows)' points to the data rows. The table itself has a light blue header row and white data rows. The data rows contain the following information:

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000

Relation Property Summary

1. Each relation name is unique
 - No two relations have the same name
2. Each cell of the relation (value of a domain) contains exactly one atomic (single) value and cannot be empty... in practice SQL allows NULL
3. Each attribute of a relation has a distinct name
4. Values of an attribute are all from the same domain
5. Each tuple is distinct. There are no duplicate tuples
 - Theoretically... in practice, SQL supports “bags” (allow duplicates)
6. Order of attributes is not important
 - Note difference from mathematical def of relations
 - Tuple (x,y) is not the same as (y,x) in mathematical definition
 - Reason: attribute names represent domain and can be reordered
7. Order of tuples is not important

Relational Model
Definitions

**Constraints and
Relationships**

Lab!

onwards...

Constraints

Relation scheme defines the types and domain of all attributes

- Can enforce constraints whenever tuples are added/modified

This can enforce many constraints to protect the integrity of your data

- Can't insert a string into an Integer type attribute
- A State field could limit domain to (AL, AK, AZ...WY)
- An SSN attribute must follow form (xxx-xx-xxxx)
- Price must be > 0.00

... but not all!

- Application or “business logic” may not be feasible
- Example: “An employee can’t work more than 40 hours per week across all jobs”

Keys

Superkey of R:

- A set of attributes that is sufficient to uniquely identify each tuple in $r(R)$

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

What is a superkey for this relation?

The *professor* relation

Keys

Superkey of R:

- A set of attributes that is sufficient to uniquely identify each tuple in $r(R)$

What is a superkey for this relation?

<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>	<i>building</i>	<i>room_number</i>	<i>time_slot_id</i>
BIO-101	1	Summer	2017	Painter	514	B
BIO-301	1	Summer	2018	Painter	514	A
CS-101	1	Fall	2017	Packard	101	H
CS-101	1	Spring	2018	Packard	101	F
CS-190	1	Spring	2017	Taylor	3128	E
CS-190	2	Spring	2017	Taylor	3128	A
CS-315	1	Spring	2018	Watson	120	D
CS-319	1	Spring	2018	Watson	100	B
CS-319	2	Spring	2018	Taylor	3128	C
CS-347	1	Fall	2017	Taylor	3128	A
EE-181	1	Spring	2017	Taylor	3128	C
FIN-201	1	Spring	2018	Packard	101	B
HIS-351	1	Spring	2018	Painter	514	C
MU-199	1	Spring	2018	Packard	101	D
PHY-101	1	Fall	2017	Watson	100	A

The *section* relation

Candidate and Primary Keys

Superkey of R:

- A set of attributes that is sufficient to uniquely identify each tuple in $r(R)$

Candidate Key of R: A “minimal” superkey

- A Candidate Key is a superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey (does not possess the superkey uniqueness property)
- A Candidate Key is a Superkey but opposite may not be true

Primary Key: The Candidate Key chosen to represent a relation/table

Super vs Candidate Key

Possible superkeys:

- $\langle \text{ID}, \text{name} \rangle$,
- $\langle \text{ID}, \text{dept_name} \rangle$,
- $\langle \text{ID}, \text{name}, \text{dept_name}, \text{salary} \rangle$

Candidate Key must be minimal:

- $\langle \text{ID} \rangle$
- $\langle \text{course_id}, \text{sec_id}, \text{semester}, \text{year} \rangle$

Primary keys are listed first and underlined when showing the schema

classroom(building, room_number, capacity)

department(dept_name, building, budget)

course(course_id, title, dept_name, credits)

instructor(ID, name, dept_name, salary)

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>	<i>building</i>	<i>room_number</i>	<i>time_slot_id</i>
BIO-101	1	Summer	2017	Painter	514	B
BIO-301	1	Summer	2018	Painter	514	A
CS-101	1	Fall	2017	Packard	101	H
CS-101	1	Spring	2018	Packard	101	F
CS-190	1	Spring	2017	Taylor	3128	E
CS-190	2	Spring	2017	Taylor	3128	A
CS-315	1	Spring	2018	Watson	120	D
CS-319	1	Spring	2018	Watson	100	B
CS-319	2	Spring	2018	Taylor	3128	C
CS-347	1	Fall	2017	Taylor	3128	A
EE-181	1	Spring	2017	Taylor	3128	C
FIN-201	1	Spring	2018	Packard	101	B
HIS-351	1	Spring	2018	Painter	514	C
MU-199	1	Spring	2018	Packard	101	D
PHY-101	1	Fall	2017	Watson	100	A

Picking a Primary Key

Every Relation must have a Primary Key

How to pick from the candidates?

- Based on business logic
- Is “Name” unique? depends on your business/application!
- Ideally Primary Key should be something that never/rarely changes

Why?

Primary Key is another type of **constraint**

- DB will enforce uniqueness of the Primary Key attributes

The magic of Databases

A database helps us **connect** multiple Relations

STUDENT

Name	Student_number	Class	Major
------	----------------	-------	-------

COURSE

Course_name	Course_number	Credit_hours	Department
-------------	---------------	--------------	------------

PREREQUISITE

Course_number	Prerequisite_number
---------------	---------------------

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
--------------------	---------------	----------	------	------------

GRADE_REPORT

Student_number	Section_identifier	Grade
----------------	--------------------	-------

How are these Relations connected to each other?

Foreign Keys

Defines a relationship connecting tuples in two relations

- The **referencing relation** and the **referenced relation**
- Defines another type of constraint - **Referential Integrity**
- **Foreign Key constraints** must connect to the Primary Key in the referenced relation

GRADE_REPORT.Student_number must match a value in **STUDENT.Student_number**

PREREQUISITE.Course_number and **Prerequisite_number** must match value in **COURSE.Course_number**, etc

etc

STUDENT

Name	Student_number	Class	Major
------	----------------	-------	-------

COURSE

Course_name	Course_number	Credit_hours	Department
-------------	---------------	--------------	------------

PREREQUISITE

Course_number	Prerequisite_number
---------------	---------------------

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
--------------------	---------------	----------	------	------------

GRADE_REPORT

Student_number	Section_identifier	Grade
----------------	--------------------	-------

Referential Integrity

Only students listed in the Students relation should be allowed to enroll for courses.

- If a value of sid appears in Enrollment relation then it **MUST** appear in Student relation
- “Only students can take courses”
- Database is automatically enforcing application requirements for you... can your Array do that?

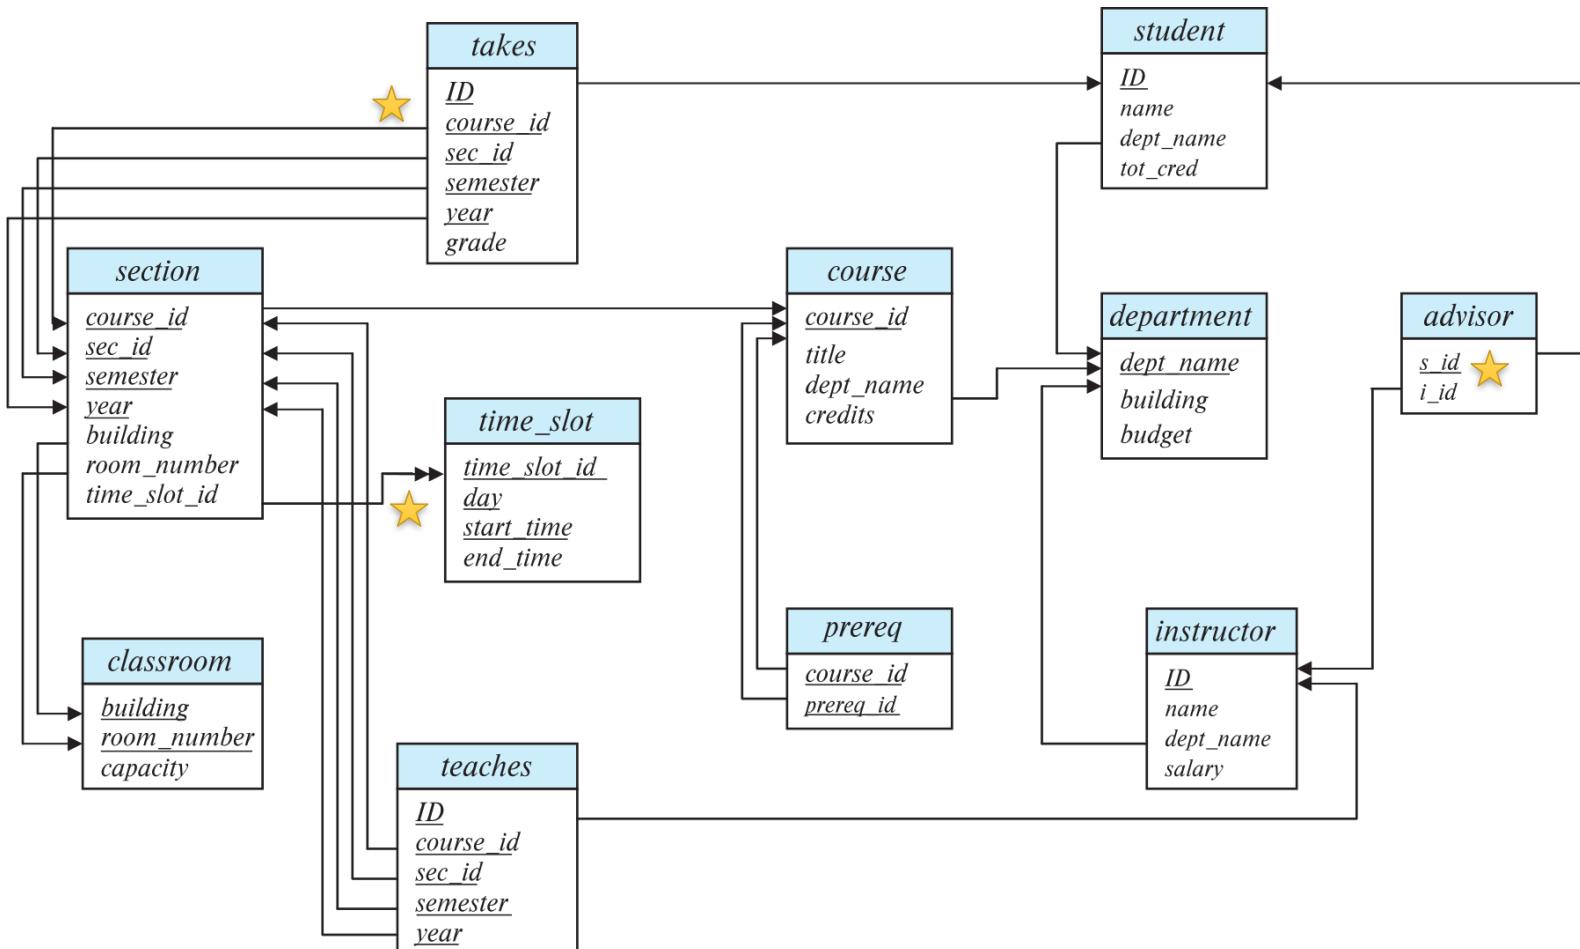
Enrollment

sid	cid	grade
53666	Jazz101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Student

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Full University Schema Diagram

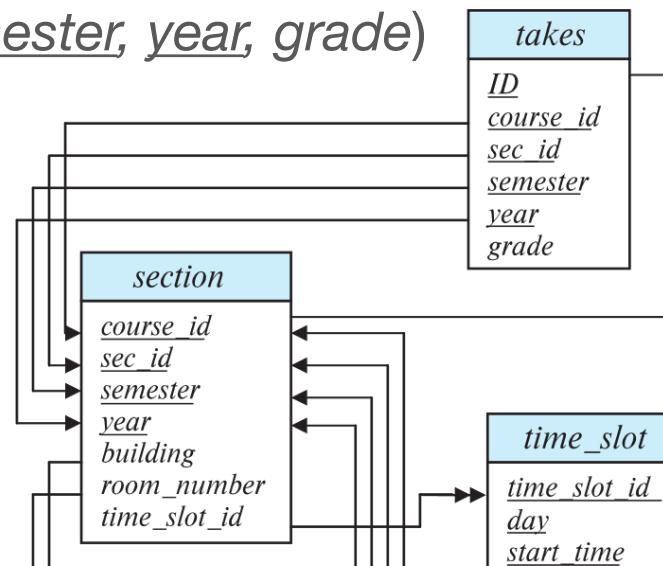


Primary Key Selection

Why do we use multiple attributes in a Primary Key?

- section(course_id, sec_id, semester, year, building, ...)
- takes(ID, course_id, sec_id, semester, year, grade)

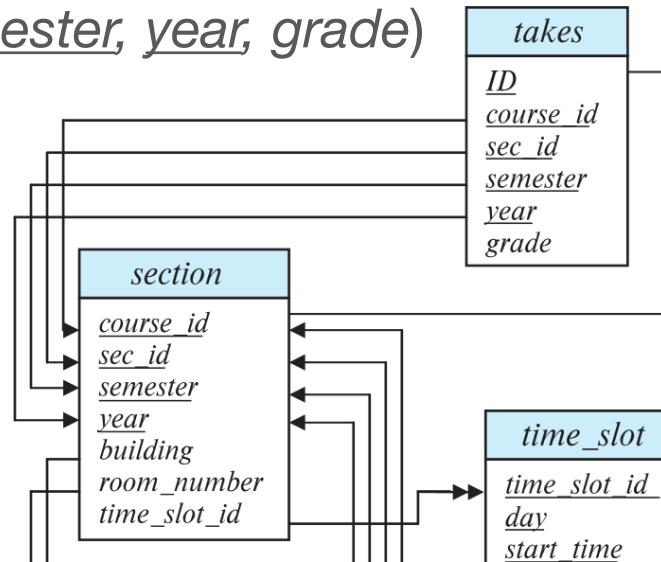
Why not just define a unique attribute like **sec_id_number** and use that by itself?



Primary Key Selection

Why do we use multiple attributes in a Primary Key?

- section(course_id, sec_id, semester, year, building, ...)
- takes(ID, course_id, sec_id, semester, year, grade)
- Using a single field looks simpler, but it prevents the benefit of the DB enforcing uniqueness
- Using sec_id_number as foreign key requires us to look up info from multiple tables which may be less efficient

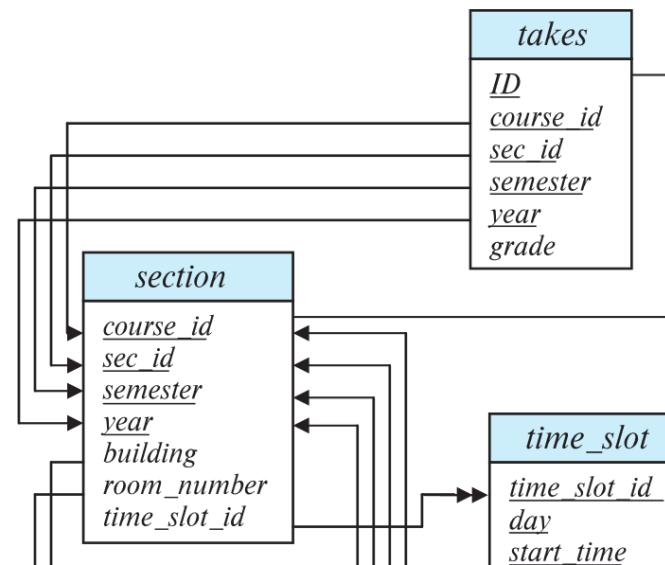


Primary Key Selection

Consider this:

- `takes(ID, course_id, sec_id, semester, year, grade)`

Does this match the
“business logic” we actually
want for our university?
(Hint: what uniqueness will
this enforce?)



Primary Key Selection

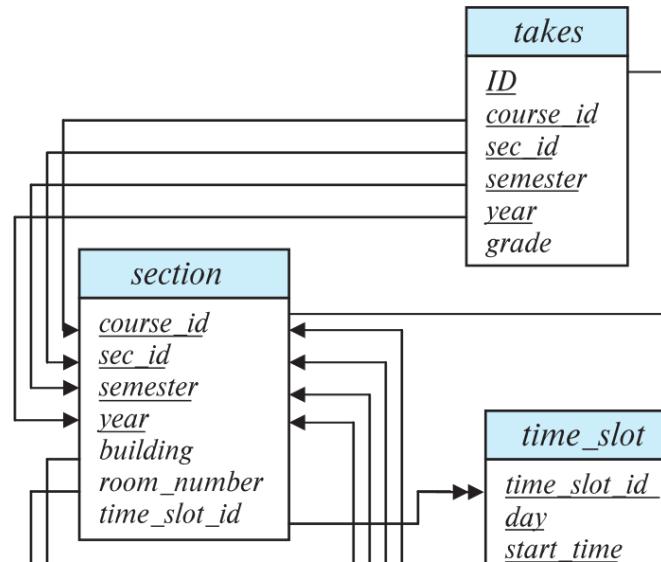
Consider this:

- takes(ID, course_id, sec_id, semester, year, grade)

This Primary Key allows a student to be registered for multiple sections of the same course at once!

But if we remove sec_id, then we will not have a complete Foreign Key!

- We must match all fields in the other relation's PK to qualify as a Foreign Key
- In practice, many SQL DBs don't support Referential Integrity without a complete Foreign Key



Relational Model
Definitions

Constraints and
Relationships

Lab!

onwards...

Lab 2: Python Flask

GW CS 2541: Database Systems and Team Projects - 2022
Prof. Tim Wood, Ethan Baron, and Catherine Meadows

Front End vs Backend: What is the difference?

Front End vs Back End

Front End:

- The art and design of websites and web applications that render on the client side. Everything from the look and feel to the way you interact with a website.

Back End:

- The server side logic for an application controlling what happens with the data, how the client side rendering changes in response and how the data gets stored.

Full Stack:

- A developer that can work on both Front and Back End software

Python Flask

Flask is a Python “Web framework”

How is this different from
a “web server”?

- Library to make it easier to write a web application
- Examples: Flask, Django, react.js, vue.js, Angular, Ruby on Rails, Drupal...

API for defining backend services

- Handles form input/output, cookie management, session data, DB connections, overall page formatting, etc

Flask is a microframework

- Provides a minimal set of functionality (with extras as needed)

Flask Hello World

Load the Flask library
and setup your app

Define a **Route** and
specify what is returned
when we access it

Tell the web application
to actually start running

```
from flask import Flask  
  
app = Flask('app')  
  
@app.route('/')  
def hello_world():  
    return 'Hello, World!'  
  
app.run(host='0.0.0.0', port=8080)
```

*Be careful about
copy/paste from
slides!*

Flask on Replit.com

The screenshot shows the Replit IDE interface. On the left, the code editor displays `main.py` with the following content:

```
1 from flask import Flask
2 app = Flask('app')
3
4 @app.route('/')
5 def hello_world():
6     return 'Hello, World!'
7
8 app.run(host='0.0.0.0', port=8080)
9
```

To the right of the code editor is a preview window showing the output of the application: "Hello, World!". Below the preview is a terminal window showing log messages:

```
* Running on all addresses.
WARNING: This is a development server. Do not
use it in a production deployment.
* Running on http://172.18.0.61:8080/ (Press CT
RL+C to quit)
172.18.0.1 - - [18/Jan/2022 15:17:28] "GET /
P/1.1" 200 -
```

Three callout bubbles provide additional information:

- A blue callout bubble points to the code editor with the text "Write code".
- A blue callout bubble points to the preview window with the text "App preview (best to open in new tab)".
- A blue callout bubble points to the terminal window with the text "Log messages - print() calls will go here".

Routes + Templates = Flask

Routes

- A backend service endpoint URL
- Function to be called when route is accessed

Templates

- Defines front end appearance of website
- Interacts with back end to allow data to be filled in

Both are good examples of our goal of Abstraction!

- Flask helps us separate design of different backend services and cleanly separates front end from back end

(Other stuff too, but this will get you most of the way to a good app!)

Routes Basics

Routes are functions:

- Specify the URL to access them
- Define the behavior to execute
- Return the content that should be displayed to the user

`@app.route` is a Python Decorator

- Special syntax to make a wrapper function. See [1] for details

```
from flask import Flask  
  
app = Flask('app')  
  
@app.route('/two')  
  
def hello_world():  
  
    Return '<html><body>This route has valid  
HTML, but both will display!</body></html>'  
  
@app.route('/')  
  
def hello_world():  
  
    return 'Hello, World from the root route!'  
  
app.run(host='0.0.0.0', port=8080)
```

[1] <https://realpython.com/primer-on-python-decorators/>

Templating Basics: Passing Variables

/templates/index.html

Templates
MUST be in a
folder
/templates

```
<html>
    <head>
        <title> {{ title }} </title>
    </head>
    <body>
        <h1> Hello {{ username }}</h1>
    </body>
</html>
```

```
from flask import Flask
from flask import render_template

app = Flask('app')

@app.route('/index')
def index():
    name = 'Cat Meadows'
    return render_template('index.html',
                          title = 'Welcome',username=name)

app.run(host='0.0.0.0', port=8080)
```

Import module for
rendering
templates

Demo: Hello X and Hello Y

Reuse:

- Template can be used in multiple routes
- Each route can fill different data into the template

```
@app.route('/')
def hello_world():
    name = 'Cat Meadows'
    return render_template('hello.html', title =
'Welcome', username=name)

@app.route('/helloTim')
def hello_world():
    name = 'Tim Wood'
    return render_template('hello.html', title = 'Welcome
2', username=name)
```

Demo:

<https://replit.com/team/cs2541s22/Lab2-Practice-Template-Data>

Routes with Parameters

We can extract data from the URL

- Parameters are available as python variables
- Flask lets you enforce types, have multiple parameters, etc

Modify the route definition and add the parameters as arguments to your function

What would we see if we visit /parameters/Ethan ?

```
from flask import Flask  
app = Flask('app')  
  
@app.route('/parameters/<name>')  
def hello_name(name):  
    return 'Hello, ' + name  
  
app.run(host='0.0.0.0', port=8080)
```

Activity 1: Hello/XYZ

Our earlier code was dumb!

- Repetitive routes that are just different based on the incoming data

Group Task:

- Make a single route which can say "hello **XYZ**" based on the URL data
- Must use template

5 minutes!

```
@app.route('/')
def hello_world():
    name = 'Cat Meadows'
    return render_template('hello.html', title =
'Welcome', username=name)

@app.route('/helloTim')
def hello_world():
    name = 'Tim Wood'
    return render_template('hello.html', title = 'Welcome
2', username=name)
```

/hello/Tim -> "Hello Tim!"
/hello/Ethan -> "Hello Ethan!"
etc...

Templating Basics: if

```
<html>  
  <head>  
    <title> {{ title }} </title>  
  </head>  
  <body>  
    {% if username == "Cat Meadows": %}  
      <h1> Hello, UTA!! </h1>  
    {% else %}  
      <h1> Hello, {{ username }} </h1>  
    {% endif %}  
  </body>  
</html>
```

Syntax {%- %}

Always close
conditionals with
{%- %}

Templating Basics: for loop

```
<html>
  <head>
    <title> {{ title }} </title>
  </head>
  <body>
    <ul>
      {% for user in users: %}
        <li> {{ user }} </li>
      {% endfor %}
    </ul>
  </body>
</html>
```

Always close for
loops with
{% endfor %}

```
from flask import Flask
from flask import render_template
app = Flask('app')

@app.route('/index')
def index():
    users = ['Tim', 'Ethan', 'Cat']
    return render_template('index.html',
                          title = 'Welcome', users=users)

app.run(host='0.0.0.0', port=8080)
```

More Template Syntax

Demo: "Lab-2-Practice-Template-Syntax" repl.it

Learn more:

- <https://realpython.com/primer-on-jinja-templating/>
- <https://blog.miguelgrinberg.com/post/the-flask-mega-tutorial-part-ii-templates>

Activity 2: Class Roster

1. Create a class roster **nested dictionary** in the format:

```
Tim = {  
    "Name": "Tim Wood"  
    "SID": "G12345678"  
    "Engagement": 2  
}  
#... define more students here ...  
roster = {  
    "Tim" : Tim,  
    "Ethan" : Ethan,  
    "Cat" : Cat  
}
```

2. Use templates and routes with parameters to display:

- First route: list of students by name at the "/" index route
 - Each name should be a link to the second route
- Second route: display name, SID, and engagement points for the student

Note: You should only use **two** templates and **two** routes

Due TOMORROW 11:59PM

<https://replit.com/team/cs2541s22/Lab2-Advanced-Routes-and-Templates>

Formatting Examples

Emphasis box 1

Emphasis box 2

Emphasis box 3

Emphasis box 4

```
from flask import Flask  
  
app = Flask('app')  
  
@app.route('/')  
def hello_world():  
    return 'Hello, World!'  
  
app.run(host='0.0.0.0', port=8080)
```